

Redis

- Redis is an open source, in-memory data structure store, used as a database.
- Stores data in form of Key Value pair
- Open Source, NoSql DB, Written in C
- Remote Dictionary Server
- strings, hashes, lists, sets, sorted sets

In-Memory Database

- Keeps data in RAM
- Means that for every interaction with database, you will only access the Main memory
- Hence the operation will be faster as it directly access main memoery instead of secondary memory

How does Redis work

All Redis data resides in-memory, in contrast to databases that store data on disk or SSDs. By eliminating the need to access disks, in-memory data stores such as Redis avoid seek time delays and can access data in microseconds. Redis features versatile data structures, high availability, geospatial, Lua scripting, transactions, on-disk persistence, and cluster support making it simpler to build real-time internet scale apps.

Redis Vs Memcached

Both Redis and MemCached are in-memory, open-source data stores. Memcached, a high-performance distributed memory cache service, is designed for simplicity while Redis offers a rich set of features that make it effective for a wide range of use cases. They work with relational or key-value databases to improve performance, such as MySQL, Postgres, Aurora, Oracle, SQL Server, DynamoDB, and more.

Redis Features

- Simple to install setup and manage
- Redis is very fast (1m qps)
- Redis is very fast because it loads the data from main memory
- Redis supports a variety of languages
- Redis supports master to slave replication

Redis Drawbacks

- It is single threaded
- It has got limited client support for consistent hashing
- It has significant overhead for persistence

Redis vs RDBMS

Redis	RDBMS
NoSql	Sql
Key value structure	Table Structure
Fast	Slow
Stores in main memory	Stores in secondary memory
Generally stores small data	Stores any kind of data
Not follow ACID	Follow ACID

AOF(Append-only files)

- To avoid data loss in case of system crash Redis provides persistence by Snapshotting & AOF(Append-only files)
- AOF copies incoming write command to secondary memory
 - always — Every writes to redis, writes to secondary memory
 - everysec — Writes to disk every second

Redis String

- Redis Strings are binary safe
- Redis string can contain any kind of data
- A String value can be at max 512 MB in length

SET key value

GET key

SET key value NX | XX

NX => set a value to the key if the key is not exist

XX => set a value to the key only if the key already exists

Redis List

- Redis Lists are simply lists of strings
- Lists are sorted by insertion order
- The max length of a list is $2^{32} - 1$ elements
- List Commands
 - LPUSH — Inserts a new element on the head
 - *LPUSH A 1*
 - RPUSH — Inserts a new element on the tail
 - *RPUSH A 2*
 - LPOP — Removes an element from the head
 - RPOP — Removes an element from tail
 - LRANGE — Prints an element from the specified position
 - *LRANGE A 0 -1*

Redis Set

- Redis Sets are an unordered collection of Strings.
- It is possible to add, remove, and get in $O(1)$
- The max number of members in a set is $2^{32} - 1$
- Commands
 - SADD — Adds a new element to the SET
 - SMEMBERS — To retrieve values from the KEY
 - SISMEMBER — To verify whether specified element is present in the SET
 - SPOP — To pop the element out from the SET
 - SUNION — Union of all the given sets.
 - SINTER — Intersection of all the given sets.
 - SUNIONSTORE — This command is equal to SUNION, but instead of returning the resulting set, it is stored in destination.
 - SINTERSTORE — This command is equal to SINTER, but instead of returning the resulting set, it is stored in destination.

Redis Hashes

- Redis Hashes are maps between string fields and string values, so they are the perfect data type to represent objects like user-details
- Every hash can store up to $2^{32} - 1$ field-value pairs
- Commands
 - HMSET — Sets the specified fields to their respective values in the hash stored at key. *Time complexity: $O(N)$*
 - HGET — Returns the value associated with field in the hash stored at key. *Time complexity: $O(1)$*
 - HGETALL — Returns all fields and values of the hash stored at key. *Time complexity: $O(N)$*

Redis Sorted sets

- Redis Sorted Sets are similar to Redis Sets with the unique feature of values stored in a set. The difference is, every member of a Sorted Set is associated with a score, that is used in order to take the sorted set ordered, from the smallest to the greatest score.
- The max number of members in a set is $2^{32} - 1$
- Commands
 - ZADD — Adds all the specified members with the specified scores to the sorted set stored at key. *Time complexity:* $O(\log(N))$
 - ZREM — Removes the specified members from the sorted set stored at key. Non existing members are ignored. *Time complexity:* $O(M \cdot \log(N))$
 - ZRANGE — Returns the specified range of elements in the sorted set stored at key. *Time complexity:* $O(\log(N) + M)$
 - ZRANGEBYSCORE — Returns all the elements in the sorted set at key with a score between min and max. *Time complexity:* $O(\log(N) + M)$

References

- <https://aws.amazon.com/redis/>
- <https://medium.com/@cosmicconvallis/interview-questions-on-redis-for-developers-c27769b410e9>